

COMP90055 Computing Project

Parallelising Brzozowski's DFA Construction Algorithm

Joshua Clark


Semester 1, 2019

Research aims, put simply...

We want to take something like this:

$$(a + b)^* \& \neg(aa)$$

And produce something like this:



And see what speedup we can get by using multiple processors.

Research aims

- What performance gains, if any, are to be had from a parallel compared to a sequential implementation?
- What advantages and/or disadvantages do imperative and functional approaches to this problem offer respectively?
- Brzozowski's algorithm depends on a notion of regular expression equivalence in order to produce minimal or near-minimal DFAs. What influence does this, and other such parameters, have on attempts at parallelisation?

(Regular) languages

We consider a language, L , to be a set of strings over some finite alphabet Σ .

$$L \subseteq \Sigma^*$$

We are concerned with *regular* languages. A language is regular if it can be described by a *regular expression* or, equivalently, if there is some *Deterministic Finite Automaton* which recognises the language.

Regular expressions

\emptyset , ε and $a \in \Sigma$ are all regular expressions.

If R and S are regular expressions then so are:

R^*

$R \cdot S$

$R + S$

$R \& S$

$\neg R$

Regular expressions

\emptyset , ε and $a \in \Sigma$ are all regular expressions.

If R and S are regular expressions then so are:

R^*

$R \cdot S$

$R + S$

$R \text{ & } S$

$\neg R$

The inclusion of the highlighted cases, while still regular operations, gives what is known as *extended* regular expressions.

Derivatives of a language

Formally, the derivative with respect to a , where $a \in \Sigma, s \in \Sigma^*$, D_a is defined as:

$$D_a L = \{s \mid as \in L\}$$

Informally, we can think of the derivative of a regular language w.r.t. symbol $a \in \Sigma$ as the strings in the language starting with that symbol, but with that symbol taken away.

$$L = \{\textcolor{red}{aa}, \textcolor{red}{ab}, \textcolor{red}{aab}, \textcolor{red}{bb}\}$$

$$D_a L = \{a, b, ab\}$$

Defining regular expression derivatives

$$D_a \emptyset = \emptyset$$

$$D_a \varepsilon = \emptyset$$

$$D_a b = \begin{cases} \varepsilon, & \text{if } a = b \\ \emptyset, & \text{otherwise} \end{cases}$$

$$D_a(R^*) = D_a R \cdot R^*$$

$$D_a(R \cdot S) = D_a R \cdot S + \mathcal{N}(R) \cdot D_a S$$

$$D_a(R + S) = D_a R + D_a S$$

$$D_a(R \& S) = D_a R \& D_a S$$

$$D_a(\neg R) = \neg D_a R$$

where

$$\mathcal{N}(R) = \begin{cases} \varepsilon, & \text{if } \text{nullable}(R) \\ \emptyset, & \text{otherwise} \end{cases}$$

We say that a set which contains the empty string is *nullable*.

Brzozowski's algorithm

```
function BUILDDFA( $q_0, \Sigma$ )
     $TODO \leftarrow \{q_0\}$ ,  $Q \leftarrow \{q_0\}$ ,  $\delta \leftarrow \emptyset$ ,  $F \leftarrow \emptyset$ 
    if  $\mathcal{N}(q_0) = \varepsilon$  then
         $F \leftarrow \{q_0\}$ 
    while  $TODO \neq \emptyset$  do
        let  $q \in TODO$ 
         $TODO \leftarrow TODO \setminus \{q\}$ 
         $Q \leftarrow Q \cup \{q\}$ 
        for  $a \in \Sigma$  do
             $q' \leftarrow D_a q$ 
             $\delta \leftarrow \delta \cup \{(q, a), q'\}\}$ 
            if  $q' \notin Q \cup TODO$  then
                 $TODO \leftarrow TODO \cup \{q'\}$ 
            if  $\mathcal{N}(q') = \varepsilon$  then
                 $F \leftarrow F \cup \{q'\}$ 
return  $(Q, \Sigma, \delta, q_0, F)$ 
```

Related work: Strauss et al. (2006, 2014, 2017)

- The authors offer a series of process-oriented specifications of Brzozowski's DFA construction algorithm in the style of Hoare's CSP
- They have implemented one of their concurrent specifications in the language Go
- Benchmarked their implementation against a test suite of randomly generated regular expressions (of varying complexity).

Parallelism and Concurrency in Haskell

As a pure language, Haskell functions produce no side-effects unless explicitly indicated. This means that the distinction drawn between parallelism and concurrency is somewhat different from that offered by Go.

- Concurrency is a method for structuring programs and relates naturally to situations involving I/O and non-determinism
- Parallelism, on the other hand, means speeding up the execution of a (pure, deterministic) program by running it on multiple processors

Evaluation strategies:


Control.Parallel.Strategies

- A Strategy is a higher level abstraction built upon Haskell's atoms for parallelism, *sparks*.
- It is essentially a description of how components of a data structure can be evaluated in parallel.
- Deterministic: the program will always return the same result (regardless of number of processors)
- Modular: decouples an algorithm from the method used to parallelise it

Dataflow parallelism: Control.Monad.Par

The idea behind the Par monad is that the programmer can specify the program in terms of a data flow model.

```
runPar $ do
    i <- new
    j <- new
    fork (put i (f 10))
    fork (put j (g 10))
    a <- get i
    b <- get j
    return (a,b)
```



Example code

-- *sequential*

```
derivatives :: RE -> [Symbol] -> [RE]
derivatives re abc = map (derive re) abc
```

-- *evaluation strategies*

```
derivatives :: RE -> [Symbol] -> [RE]
derivatives re abc = map (derive re) abc
                           'using' parList rdeepseq
```

-- *the Par monad*

```
derivatives :: RE -> [Symbol] -> Par [RE]
derivatives re abc = parMapM (return . derive re) abc
```

Experimental setup

The basic approach follows that employed by Strauss et al. in the hope of allowing for meaningful comparisons to be drawn between the results obtained.

- generate two sets of regular expressions (with different depth syntax trees), with $|\Sigma| = 4$ and $|\Sigma| = 94$ respectively.
- compare running times for the different implementations: sequential, strategies (with and without “chunking”) and the Par monad.
- speedup and efficiency compared to number of cores utilised are considered as metrics of overall performance

All code was compiled using GHC version 8.6.5 and run on a 2015 iMac with a 3.1GHz 4 core Intel i5 processor and 8GB of RAM.

Results: Speedup, $|\Sigma| = 4$




Figure: Speedup versus number of cores, $|\Sigma| = 4$

Results: Efficiency, $|\Sigma| = 4$




Figure: Efficiency versus number of cores, $|\Sigma| = 4$

Results: Speedup, $|\Sigma| = 94$




Figure: Speedup versus number of cores, $|\Sigma| = 94$

Results: Efficiency, $|\Sigma| = 94$




Figure: Efficiency versus number of cores, $|\Sigma| = 94$

Conclusions

- The Par monad approach outperformed the Evaluation Strategies approach for both the small and larger alphabet.
- Generally the boost in speedup to be gained and overall efficiency diminished with an increased number of cores.
- More significant performance gains were achieved when the alphabet was larger.
 - interestingly, this reverses the findings reported by Strauss et al. (they achieved a median speedup of 1.72 for the small alphabet, yet only 1.09 for the large alphabet case)